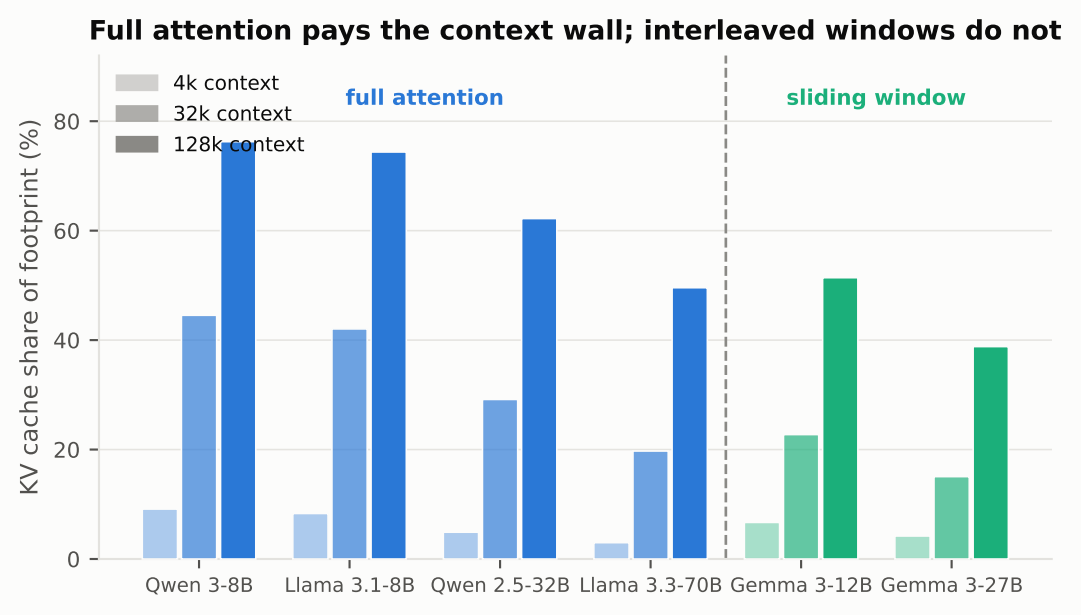
Full attention pays the context wall; interleaved windows do not
4k context
32k context
128k context
full attention
sliding window
KV cache share of footprint (%)
0
20
40
60
80
Qwen 3-8B
Llama 3.1-8B
Qwen 2.5-32B
Llama 3.3-70B
Gemma 3-12B
Gemma 3-27B